

An introductory course on General Purpose Computing on GPUs

Additional details about the exercises

Francesco Lettich

Insight Lab

Universidade Federal do Ceará, Fortaleza, Brasil



Introduction

- ✦ To complete the exercises you will **need**:
 - ✦ Recommended O/S: Linux (e.g., Ubuntu \geq 16.04).
 - ✦ CUDA 9.
 - ✦ A GCC version compatible with CUDA 9 (GCC v5.x or v6.x are ok).
 - ✦ A NVIDIA videocard.
 - ✦ CMake.
- ✦ I will also give you a zip containing the **source code** of a little application that you will have to **complete** by providing the **required solutions**.

Setup

- ✦ To download CUDA 9:

https://developer.nvidia.com/cuda-downloads?target_os=Linux

- ✦ When installing CUDA 9, please follow the instructions provided by NVIDIA, especially the modifications you have to apply to the “.bashrc” file; more info here (**Section 4.1.5**):

<http://docs.nvidia.com/cuda/cuda-quick-start-guide/index.html>

- ✦ To download **GCC 5.x or 6.x**, please use “*apt install*” with gcc-5/g++-5 or gcc-6/g++-6 (or something equivalent).

- ✦ To install CMake, type “*apt install cmake*”.

Archive details

- ✦ Within the archive you will find the following files/folders:
 - ✦ **CmakeLists.txt** => file used by the tool “Cmake” to “prepare” the compilation (generates some files).
 - ✦ **main.cu** => The file containing the “main” (entry point) of the program. Notice the extension “**cu**”, which tells the nvcc compiler that main.cu contains C++ code augmented with CUDA keywords and syntax.
 - ✦ **kernels.h** and **kernels.cu** => respectively containing the **declarations** and the **definitions** of the kernels (functions) you have to provide as part of the **exercises**.

Compilation and execution

- ✦ Assuming that you have a computer **properly configured**, go in the folder where you unzipped the archive...then:
 - ✦ To “**prepare**” the **compilation**, type “*cmake CmakeLists.txt*”.
 - ✦ Then, to **compile** the source code type “*make*”.
 - ✦ Finally, to **execute** the binary go in the “/bin” subfolder and execute the “*main*” executable.

List of Exercises / 1

- ✦ **Exercise 1**: implement the GPU-based version of the **Hillis-Steele** algorithm, limitedly to the case of a **single thread-block**. This is represented by the function “*HillisAndSteele*” in kernels.cu.
- ✦ **Exercise 2**: implement the Blelloch’s (down/up-sweep) algorithm, limitedly to the case of a **single thread-block**. This is represented by the function “*BlellochSingleBlock*” in kernels.cu.
- ✦ **Exercise 3**: Adapt the code written for exercise 2 to implement the **inclusive prefix sum** (easy!). This is represented by the function “*BlellochSingleBlockInclusive*” in kernels.cu.

List of Exercises / 2

- ✦ Exercise 4: Extend the code of **exercise 2** to make use of **multiple thread-blocks**. This is represented by the function “*BlellochMultipleBlocks*” in kernels.cu.
- ✦ Exercise 5 (optional!): Modify the source code of the solution of exercise 2 to avoid **bank conflicts** in shared memory. This is represented by the function “*BlellochNoConflicts*” in kernels.cu.

More on the exercises

- ✦ To complete the exercises you will have to complete part of the *main function* in **main.cu** and provide the body of the functions in **kernels.cu**.
- ✦ Please, **read carefully** the **notes** I've **left** in the source code!
- ✦ Feel free to adapt the source code to your needs: the only requirement is that it should be **easy** to **verify** the **correctness** of your solutions!
- ✦ Once you complete the exercises, please **mail back** to me the zip with **your solutions** (for the e-mail, see the next slide).

Questions & Doubts

If you have relevant questions, doubts, problems, or you want to send me the solutions, please contact me at:

francesco.lettich@gmail.com